

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

# Apache Kafka

Riccardo Scotti

Università di Bologna  
Distributed Software Systems  
Prof. Paolo Ciancarini

# Agenda

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

- 1 Introduction
- 2 Context
- 3 Architectural drivers
- 4 Structure
- 5 Behavior
- 6 Rationale
- 7 Similar or competing middlewarees
- 8 Conclusions

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

# Introduction

# Introduction

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions



- Distributed platform for event streaming
- Designed to handle real-time data
- Based on pub/sub
- Originally developed at LinkedIn

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

# Context

# Use cases and scenarios

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

- What are Kafka's main use cases?

# Use cases and scenarios

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

- What are Kafka's main use cases?
- High flexibility  $\Rightarrow$  many scenarios:
  - Real-time data analytics
  - Aggregated logs
  - Data broker

# Real-time data analytics

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers


Structure

Behavior

Rationale

Similar or  
competing  
middleware

Conclusions



LinkedIn

- Activity data
- Operational metrics

DataSift

- User consumption tracker

Twitter

- Part of Twitter Storm



# Aggregated logs

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

- Exploits Kafka's distributed commit log architecture
- Already employed by Kafka's development team [6]

## Building a Replicated Logging System with Apache Kafka

Guozhang Wang<sup>1</sup>, Joel Koshy<sup>1</sup>, Sriram Subramanian<sup>1</sup>, Kartik Paramasivam<sup>1</sup>

Mammad Zadeh<sup>1</sup>, Neha Narkhede<sup>2</sup>, Jun Rao<sup>2</sup>, Jay Kreps<sup>2</sup>, Joe Stein<sup>3</sup>

<sup>1</sup>LinkedIn Corporation, <sup>2</sup>Confluent Inc., <sup>3</sup>Big Data Open Source Security LLC

### ABSTRACT

Apache Kafka is a scalable publish-subscribe messaging system with its core architecture as a distributed commit log. It was originally built at LinkedIn as its centralized event pipelining platform for online data integration tasks. Over the past years developing and operating Kafka, we extend its log-structured architecture as a replicated logging backbone for much wider application scopes in the distributed environment. In this abstract, we will talk about our design and engineering experience to replicate Kafka logs for various distributed data-driven systems at LinkedIn, including source-of-truth data storage and stream processing.

Recently, there has been much renewed interest in using log-centric architectures to build distributed systems that provides efficient durability and availability [3, 5, 6, 7]. In this approach, a collection of distributed servers can maintain a consistent system state via a replicated log that records state changes in sequential order. When some of the servers fail and come back, their states can be deterministically reconstructed by replaying this log upon recovery. At LinkedIn, the idea of using a reliable and highly available log structure as the underlying data flow to scale distributed systems has also been well envisioned and exercised over the past years. Given its log-structured architecture, we realize Kafka could be a good fit for this vision and have extended its design as a replicated logging system. In this abstract, we

# Data broker

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

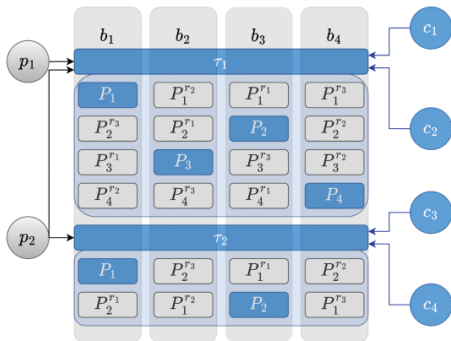
Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

- Publisher/subscriber communication paradigm
- Fine tuned data exchange policies



Apache Kafka

Riccardo  
Scotti

Introduction

Context

**Architectural  
drivers**

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

# Architectural drivers

# Architectural drivers

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

Most of them are a direct consequence of being developed for LinkedIn [1]:

- Low latency & high throughput
- Scalability
- Fault tolerance
- Consistency
- Interoperability

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

**Structure**

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

# Structure

# Components

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

The basic unit of information in Kafka is called **message**, a key-value pair

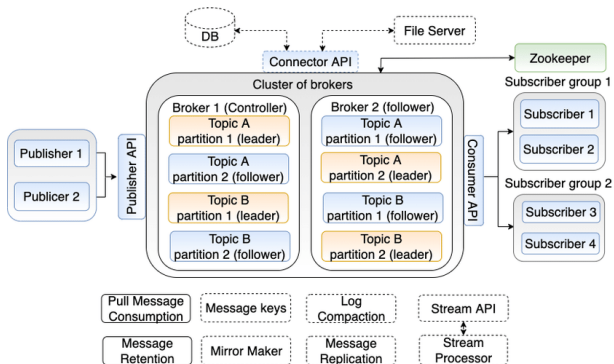
- Topics and partitions
- Producers and consumers
- Brokers and clusters
- ZooKeeper (almost completely replaced by KRaft [5])
- Connectors
  - Kafka Connect
  - Source and sink connectors
  - Connect workers

# Architectural patterns

Apache Kafka

Riccardo  
Scotti

Kafka is a pub/sub, distributed commit log [5]



Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

**Behavior**

Rationale

Similar or  
competing  
middlewarees

Conclusions

# Behavior



# Production

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

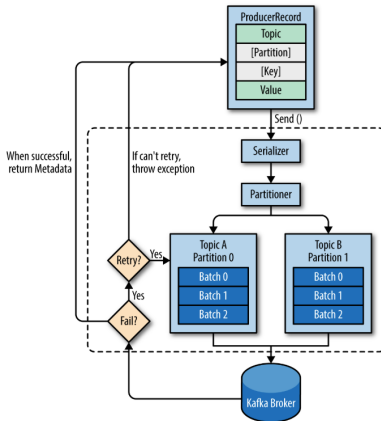
Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



# Consumption

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers


Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



# Brokers and clusters

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

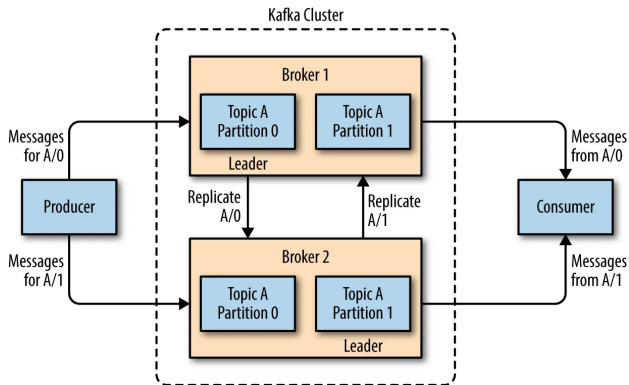
Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



# ZooKeeper

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

- Part of the Kafka environment, but a separate Apache system (introduced one year before Kafka in [2])
- Takes care of coordination
  - Storing metadata
  - Elects controller brokers
  - Assigns partitions
- No fixed role, community decides best practices
- Currently being removed, will be replaced by KRaft

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

**Rationale**

Similar or  
competing  
middlewarees

Conclusions

# Rationale

# Rationale

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions

- Publisher/subscriber model
  - Seamless communication between producers and consumer
  - Space decoupling
- Partitions and distributed commit log
  - Horizontal scalability
  - High availability and parallel reads
  - Fault tolerance
- Consumer groups
  - Handling high volumes of data
- Log-based architecture
  - Ensure ordering of events
  - Easily rebuilding state after a failure
- Retention policy
  - Data durability for aggregated logs or monitoring

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

## Similar or competing middlewares

# RabbitMQ

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



- Message broker with load balancing
- Based on point-to-point message queues
- No message retention policy
- Priority messages



# Apache Pulsar

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewares

Conclusions



- Cloud native, publisher/subscriber platform for messaging, streaming and queueing
- Similar architecture to Kafka, more complex
- BookKeeper and RocksDB to handle long-term persistence
- Less efficient
- Not as mature and widespread as Kafka

# NATS

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



- Lightweight, message-oriented middleware
- Designed for message exchange with low latency and high performances
- Minimal additional functionalities
- Well suited for IoT contexts
- Less versatile than Kafka

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

# Conclusions

# Conclusions

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions

To sum up Kafka:

- Pub/sub platform
- Distributed commit log
- Real-time data
- Analytics and brokering
- Availability, scalability and fault tolerance
- Big presence on the market (LinkedIn, Netflix, Twitter)
- Large open-source community

# References I

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



GARG, N.

*Apache kafka.*

Packt Publishing Birmingham, UK, 2013.



HUNT, P., KONAR, M., JUNQUEIRA, F. P., AND  
REED, B.

{ZooKeeper}: Wait-free coordination for internet-scale  
systems.

*In 2010 USENIX Annual Technical Conference (USENIX  
ATC 10) (2010).*



LAZIDIS, A., TSAKOS, K., AND PETRAKIS, E.

Publish-subscribe approaches for the iot and the cloud:  
Functional and performance evaluation of open-source  
systems.

*Internet of Things 19 (05 2022), 100538.*

# References II

Apache Kafka

Riccardo  
Scotti

Introduction

Context

Architectural  
drivers

Structure

Behavior

Rationale

Similar or  
competing  
middlewarees

Conclusions



RAPTIS, T. P., AND PASSARELLA, A.

A survey on networked data streaming with apache kafka.  
*IEEE Access* 11 (2023), 85333–85350.



SHAPIRA, G., PALINO, T., AND SIVARAM, R.

*Kafka - the definitive guide*, 2 ed.  
O'Reilly Media, Sebastopol, CA, Nov. 2021.



WANG, G., KOSHY, J., SUBRAMANIAN, S.,  
PARAMASIVAM, K., ZADEH, M., NARKHEDE, N., RAO,  
J., KREPS, J., AND STEIN, J.

Building a replicated logging system with apache kafka.  
*Proc. VLDB Endow.* 8, 12 (aug 2015), 1654–1655.